

# Natural Language Processing

---

Jim Martin -- Lecture 2  
CSCI 5832



# Shocking



NYTIMES.COM

## **Electric Eels Hunt in Packs, Shocking Prey and Scientists**

The behavior, used by wolves and orcas to run down fast prey, is rarely seen i...

# Today

- Words
- Corpora
- Vocabularies
- Text normalization
- Subword units

# Complex Morphology

Many languages require complex morphological segmentation/analysis

- ◆ Turkish
- ◆ Uygarlastiramadiklarimizdanmissinizcasina
- ◆ `(behaving) as if you are among those whom we could not civilize'
- ◆ Uygar `civilized' + las `become'
  - + tir `cause' + ama `not able'
  - + dik `past' + lar `plural'
  - + imiz `p1pl' + dan `abl'
  - + mis `past' + siniz `2pl' + casina `as if'

# Complex Morphology

- Morphological analysis takes a surface form and returns the stem with a set of morphological features.
  - ♦ Cats → Cat+PL   Ate → eat+PAST
- Generation goes in the other direction.
- Capturing the underlying morphology of a language is very very hard. Typically, involving a combination of finite state transducers and probabilities
- For many languages, like English and Chinese, you can get away with simpler approaches.

# HW 1 Part 1: 50 Points

- How many words do you know in your native language?
  - ◆ Due Monday 8/30 by 11:59PM
  - ◆ Submit via Canvas
  - ◆ Your answer and a writeup explaining your answer.
    - No longer than necessary. 2-3 pages should suffice.
    - Long enough to say something interesting
    - PDF; follow the naming convention specified on the Canvas assignment page

# HW 1 Part 1

- How many words do you know in your native language?
  - ◆ In your answer clearly address
    - “how many”
    - “words”
    - “know”

# Corpora and Vocabularies



# Vocabulary

- What is a vocabulary (or lexicon)?
  - ◆ All the words in a language?
    - All the words in a comprehensive dictionary?
  - ◆ Or some subset?
    - The words needed for a particular application?
      - How many?
      - How do we determine the right list?
      - Is it fixed or will we be adding new words?

# Corpora

- Words don't appear out of nowhere.
- We typically generate vocabularies from collections of representative text for a possible domain or application.
- But texts are produced by specific writer(s), at a specific times, in a variety of a specific languages, for specific functions. All of which combine combinatorially.

# Dimensions of Corpora

- Language
  - ♦ 7097 languages in the world (not all have written forms)
- Language variety
  - ♦ Dialects, creoles, pidgins, etc.
- Code switching
  - ♦ Mixing language in the same utterance
    - S/E: Por primera vez veo a @username actually being hateful! It was beautiful:)  
*[For the first time I get to see @username actually being hateful! it was beautiful:]*
    - H/E: dost tha or ra- hega ... dont worry ... but dherya rakhe  
*[“he was and will remain a friend ... don’t worry ... but have faith”]*
- Genre
  - newswire, fiction, non-fiction, scientific articles, Wikipedia
- Demographics
  - writer's age, gender, race, socioeconomic status, etc.
- Medium
  - ♦ Spoken, written, captioned, w/ video

# Corpora Metrics

$N$  = number of instances or tokens in a corpus

$V$  = vocabulary = set of unique types

$|V|$  is the size of the vocabulary

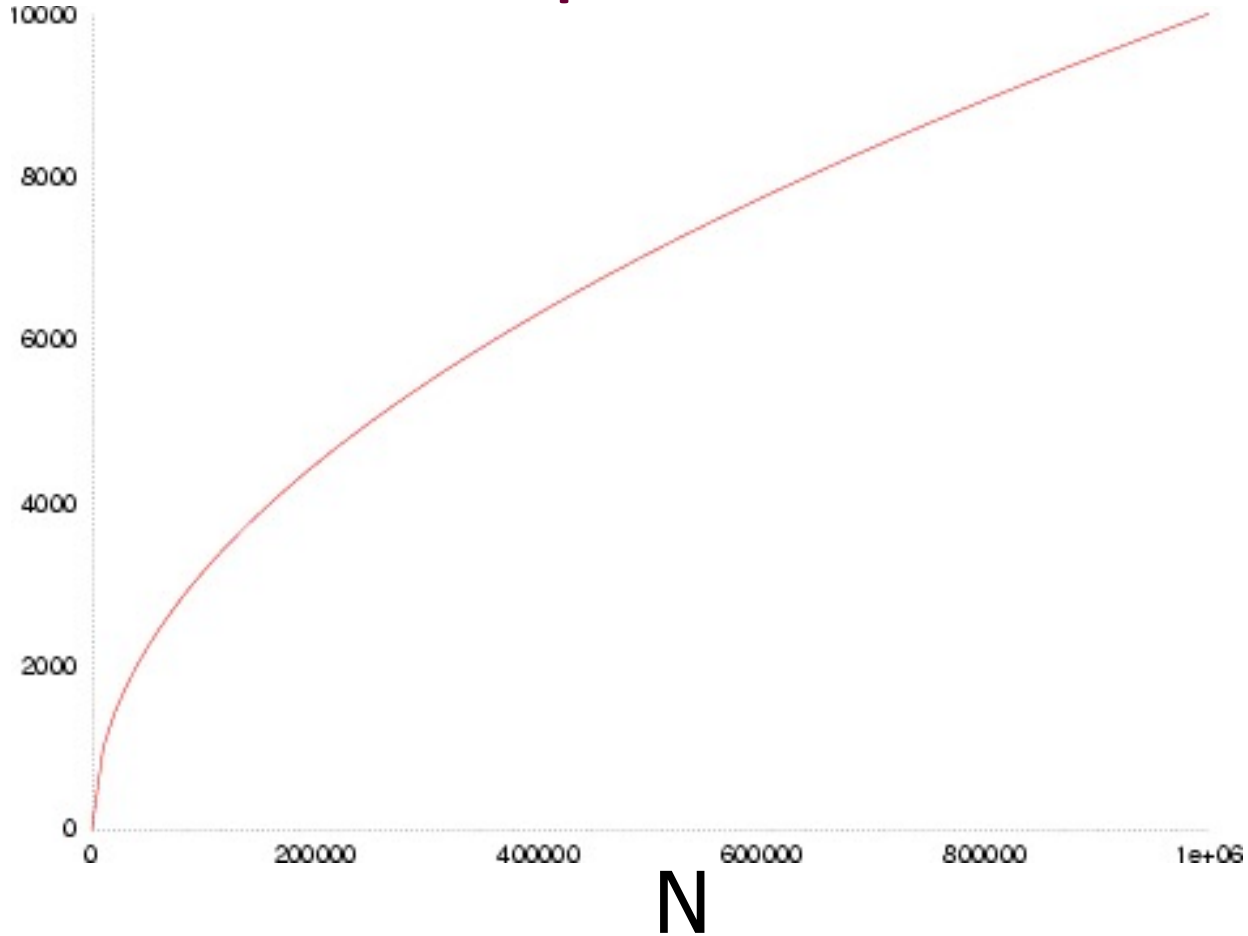
	Tokens = $N$	Types = $ V $
Switchboard corpus (phone conversations)	2.4 million	20 thousand
Shakespeare	884,000	31 thousand
Google N-grams	1 trillion	13 million

# Heaps' Law

- Heaps Law says that the size of the vocabulary grows somewhat faster than the square root of the corpus size  $|V| = kN^{\beta}$
- Where  $k$  and  $\beta$  are free variables. Usually  $k$  is between 10 and 100 and  $\beta$  is between .6 and .7
- Which really means that the rate of growth of the vocabulary tails off as the corpus grows but never completely flattens out.

# Heaps' Law

$|V|$



# Building a Vocabulary

- Let's say we're handed some large text collection representative of some application domain.
  - ◆ Like building a Shakespeare generator
- How do we pull the words out from the text to form a vocabulary?

# Simple Tokenization in UNIX

- Given a text file, output all the word types and their associated frequencies in a given text corpus
  - ◆ Inspired by Ken Church's UNIX for Poets.
- Unix has many commands to deal with basic text processing operations
  - ◆ Original Unix designers cared a lot about text processing



# Notebook

# Issues in Tokenization

Of course, it is never really that easy. There are lots of complications.

- Finland's capital → Finland Finlands Finland' s
- what're, I'm, isn't → What are, I am, is not
- Hewlett-Packard → Hewlett Packard
- state-of-the-art → state of the art
- Lowercase → lower-case lowercase lower case
- San Francisco → one token or two?
- m.p.g., PhD. → ??

# Tokenization: Language Issues

- French
  - ♦ *L'ensemble* → one token or two?
    - *L* ? *L'* ? *Le* ?
- German noun compounds
  - ♦ *Plastikwasserflaschenhalter*
  - ♦ 'plastic water bottle holder'

# Tokenization: language issues

- Chinese has no spaces between words
  - ◆ 莎拉波娃现在居住在美国东南部的佛罗里达。
  - ◆ 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
  - ◆ Sharapova now lives in US southeastern Florida
- Japanese allows intermingled alphabets

# Case Study:

## Word Segmentation in Chinese

- Chinese words are composed of characters
  - ◆ Characters are generally 1 syllable and 1 morpheme
  - ◆ Average word length is 2.4 characters

姚明进入总决赛



# Option 1

Syntax/Semantics driven segmentation

姚明进入总决赛

姚明      进入      总决赛

YaoMing reaches finals

## Option 2

More fine-grained segmentation

姚明进入总决赛

姚 明 进 入 总 决 赛

Yao Ming reaches overall finals

## Option 3

Since all the characters have meanings  
just use them.

姚明进入总决赛

姚 明 进 入 总 决 赛

Yao Ming enter enter overall decision game



# Large Vocabularies

- How do we make sure we're dealing with all the high frequency words that Zipf's law predicts
- While still providing a way to deal with out-of-vocabulary (OOV) terms
- And still keep the vocabulary size reasonable?

# Subword Tokenization

- Use **subword tokenization** to find words and common subwords empirically
  - ◆ Let the data tell us what the words are
- Can include common morphemes like *-est* or *-er*.
  - ◆ (A morpheme is the smallest meaning-bearing unit of a language; *unlikeliest* has morphemes *un-*, *likely*, and *-est*.)

# Subword Tokenization

- Three common algorithms:
  - ◆ **Byte-Pair Encoding (BPE)** (Sennrich et al., 2016)
  - ◆ **Unigram language modeling tokenization** (Kudo, 2018)
  - ◆ **WordPieces** (Schuster and Nakajima, 2012)
- All have 2 parts:
  - ◆ A token **learner** that takes a raw training corpus and induces a vocabulary.
  - ◆ A token **segmenter** that takes an input and tokenizes it according to a vocabulary.
    - Words present the vocabulary are left alone (unsegmented)
    - OOV words are broken into optimal sequences of words and subwords.

# Byte Pair Encoding (BPE)

Let initial vocabulary be the set of all individual characters

= {A, B, C, D,...,a, b, c, d...., 0-9, etc.}

- Repeat:
  - ♦ choose the two symbols that are most frequently adjacent in training corpus (say 'A', 'B'),
  - ♦ add a new merged symbol 'AB' to the vocabulary
  - ♦ replace every adjacent 'A' 'B' in corpus with 'AB'.
- Until  $k$  merges have been done.

# BPE Token Learner Algorithm

**function** BYTE-PAIR ENCODING(strings  $C$ , number of merges  $k$ ) **returns** vocab  $V$

$V \leftarrow$  all unique characters in  $C$                       # initial set of tokens is characters

**for**  $i = 1$  **to**  $k$  **do**                                              # merge tokens til  $k$  times

$t_L, t_R \leftarrow$  Most frequent pair of adjacent tokens in  $C$

$t_{NEW} \leftarrow t_L + t_R$                                               # make new token by concatenating

$V \leftarrow V + t_{NEW}$                                               # update the vocabulary

    Replace each occurrence of  $t_L, t_R$  in  $C$  with  $t_{NEW}$                       # and update the corpus

**return**  $V$

# Byte Pair Encoding (BPE)

- Most subword algorithms are run with initial white-space separated tokens.
- So first add a special end-of-word symbol '\_\_\_\_' before whitespace in training corpus
- Next, separate tokens into letters

# BPE Token Learner

## Original “corpus”

*low low low low low lowest lowest newer newer  
newer newer newer newer wider wider wider new  
new*

Add end-of-word tokens and segment:

**corpus**

5    l o w \_  
2    l o w e s t \_  
6    n e w e r \_  
3    w i d e r \_  
2    n e w \_

**vocabulary**

\_, d, e, i, l, n, o, r, s, t, w

# BPE Token Learner

Original “corpus”

*low\_*

*lowest\_*

*lowest\_*

*newer\_*

*newer\_*

*newer\_*

*wider\_*



# BPE Token Learner

Original “corpus”

*low \_*

*lowest \_*

*lowest \_*

*newer \_*

*newer \_*

*newer \_*

*wider \_*

# BPE token learner

## corpus

5    l o w \_  
2    l o w e s t \_  
6    n e w e r \_  
3    w i d e r \_  
2    n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w

# BPE token learner

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w e r \_  
3 w i d e r \_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w

Merge **e r** to **er**

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w er \_  
3 w i d er \_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er

# BPE

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w e r \_  
3 w i d e r \_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er

Merge **er \_** to **er\_**

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w e r\_  
3 w i d e r\_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er, er\_

# BPE

## corpus

5 l o w \_  
2 l o w e s t \_  
6 n e w er\_  
3 w i d er\_  
2 n e w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er, er\_

Merge **n** **e** to **ne**

## corpus

5 l o w \_  
2 l o w e s t \_  
6 ne w er\_  
3 w i d er\_  
2 ne w \_

## vocabulary

\_, d, e, i, l, n, o, r, s, t, w, er, er\_, ne

# BPE

Continuing the next merges are:

Merge	Current Vocabulary
(ne, w)	—, d, e, i, l, n, o, r, s, t, w, er, er—, ne, new
(l, o)	—, d, e, i, l, n, o, r, s, t, w, er, er—, ne, new, lo
(lo, w)	—, d, e, i, l, n, o, r, s, t, w, er, er—, ne, new, lo, low
(new, er—)	—, d, e, i, l, n, o, r, s, t, w, er, er—, ne, new, lo, low, newer—
(low, —)	—, d, e, i, l, n, o, r, s, t, w, er, er—, ne, new, lo, low, newer—, low—

# BPE token segmentation algorithm

- On test data
  - ♦ White-space separate
  - ♦ If character sequence is in the vocab just leave it alone
  - ♦ If its OOV
    - Run each merge learned from the training data:
    - Greedily
    - In the order we learned them
- So: merge every **e r** to **er**, then merge **er \_** to **er\_**, etc.
- Result:
  - ♦ Test set "n e w e r \_" would be tokenized as a full word
  - ♦ Test set "l o w e r \_" would be two tokens: "low er\_"

# BPE

- With BPE (and other subword approaches) there are no out of vocabulary words. Every word can be decomposed into a sequence of known vocabulary items (sequences of words and subwords, or worst case, characters).



# BERT Vocabulary

- All modern language models (BERT, GPT, T5 and their variants) and MT systems make use of relatively small vocabularies derived using one of the popular subword unit algorithms.
- Typically reported at around 30k entries. This is largely chosen for computational efficiency reasons.
- The original BERT vocabulary was generated from the “Books” corpus and an English Wikipedia dump.

# HW 1 Part 2

- In this part of the HW you'll explore a generic BERT vocabulary, with a particular focus on that 30k size. We're interested in how many words does BERT really know?
- How does BERT's vocabulary stack up against the kind of considerations we've been discussing for people?

# Notebook

# HW 1 Part 2: 50 Points

- How many words does BERT really know?
  - ◆ Due Wednesday 9/1 by 11:59PM
  - ◆ Submit via Canvas
  - ◆ Your answer and a writeup explaining your answer.
    - No longer than necessary. 2-3 pages should suffice.
    - Long enough to say something interesting
    - PDF; follow the naming convention specified on the Canvas assignment page

# Next Time

- Chapter 3. N-Gram language modeling